

# Parent and Child Synchronization

```
int
main(void)
{
    pid_t    pid;

+   TELL_WAIT();
+
    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) {
+       WAIT_PARENT();          /* parent goes first */
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
+       TELL_CHILD(pid);
    }
    exit(0);
}

static void
charatime(char *str)
{
    char      *ptr;
    int       c;

    setbuf(stdout, NULL);          /* set unbuffered */
    for (ptr = str; (c = *ptr++) != 0; )
        putc(c, stdout);
}
```

# Parent and Child Synchronization

- To synchronize, consider
  - Child sending SIGUSR2 to the parent
  - Parent sending SIGUSR1 to the child

# Routines for Parent-Child Synchronization

```
static volatile sig_atomic_t sigflag;
static sigset_t newmask, oldmask, zeromask;

static void sig_usr(int signo)
{
    sigflag = 1;
}

void TELL_WAIT(void)
{
    if (signal(SIGUSR1, sig_usr) == SIG_ERR)
        err_sys("signal(SIGUSR1) error");
    if (signal(SIGUSR2, sig_usr) == SIG_ERR)
        err_sys("signal(SIGUSR2) error");
    sigemptyset(&zeromask);
    sigemptyset(&newmask);
    sigaddset(&newmask, SIGUSR1);
    sigaddset(&newmask, SIGUSR2);
    //Block SIGUSR1 and SIGUSR2, and save current signal mask
    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
        err_sys("SIG_BLOCK error");
}
```

```
main(void)
{
    pid_t pid;
    TELL_WAIT();
    if ((pid = fork()) == 0) {
        WAIT_PARENT();
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
        TELL_CHILD(pid);
    }
    exit(0);
}
```

# Routines Contd...

```
void WAIT_PARENT(void) //Sigsuspend to sleep while waiting for a signal to occur
{
    while (sigflag == 0) sigsuspend(&zeromask);
    sigflag = 0;
    // Reset signal mask to original value
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
        err_sys("SIG_SETMASK error");
}
```

```
void TELL_CHILD(pid_t pid)
{
    kill(pid, SIGUSR1);
}
```

```
main(void)
{
    pid_t    pid;

    TELL_WAIT();

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) {
        WAIT_PARENT();
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
        TELL_CHILD(pid);
    }
    exit(0);
}
```

# abort() Function

- Causes abnormal program termination
- `void abort(void);`  
This function never returns
- `abort()` sends SIGABRT to the process.
- ANSI C requires that if the signal is caught and the signal handler returns, `abort` still does not return to the caller
- By catching SIGABRT, the process can be allowed to perform any cleanup

# Implementation of POSIX.1 abort()

```
void abort(void) /* POSIX-style abort() function */
{
    sigset_t      mask;
    struct sigaction action;
    /* Caller can't ignore SIGABRT, if so reset to default */
    sigaction(SIGABRT, NULL, &action);
    if (action.sa_handler == SIG_IGN) {
        action.sa_handler = SIG_DFL;
        sigaction(SIGABRT, &action, NULL);
    }

    if (action.sa_handler == SIG_DFL)
        fflush(NULL); /* flush all open stdio streams */
    /* Caller can't block SIGABRT; make sure it's unblocked */
    sigfillset(&mask);
    sigdelset(&mask, SIGABRT); /* mask has only SIGABRT turned off */
    sigprocmask(SIG_SETMASK, &mask, NULL);
    kill(getpid(), SIGABRT); /* send the signal */
    /* If we're here, process caught SIGABRT and returned */
    fflush(NULL); /* flush all op
    action.sa_handler = SIG_DFL;
    sigaction(SIGABRT, &action, NULL); /* rese
    sigprocmask(SIG_SETMASK, &mask, NULL); /*
    kill(getpid(), SIGABRT); /*
    exit(1); /* this should never be execute
}
```



# system() Function

- Implementation of system without signal handling

```
int system(const char *cmdstring) {
    pid_t pid; int status;
    if (cmdstring == NULL) return(1); /* always a command processor with UNIX */
    if ((pid = fork()) < 0) status = -1; /* probably out of processes */
    else if (pid == 0) { /* child */
        execl("/bin/sh", "sh", "-c", cmdstring, (char *)0);
        _exit(127); /* execl error */
    } else { /* parent */
        while (waitpid(pid, &status, 0) < 0) {
            if (errno != EINTR) {
                status = -1; /* error other than EINTR from waitpid() */
                break;
            }
        }
    }
    return(status);
}
```

# Using system() to Invoke *ed* Editor

```
static void sig_int(int signo)
{
    printf("caught SIGINT\n");
}

static void sig_chld(int signo)
{
    printf("caught SIGCHLD\n");
}

int main(void)
{
    if (signal(SIGINT, sig_int) == SIG_ERR)
        err_sys("signal(SIGINT) error");
    if (signal(SIGCHLD, sig_chld) == SIG_ERR)
        err_sys("signal(SIGCHLD) error");
    if (system("/bin/ed") < 0)
        err_sys("system() error");
    exit(0);
}
```



# Execution

```
$ ./a.out
```

```
a append text to the editor's buffer
```

```
Here is one line of text
```

```
. period on a line by itself stops append mode
```

```
1,$p print first through last lines of buffer to see what's there
```

```
Here is one line of text
```

```
w temp.foo write the buffer to a file
```

```
25 editor says it wrote 25 bytes
```

```
q and leave the editor
```

```
caught SIGCHLD
```

# Execution

```
$ ./a.out
```

```
a          append text to the editor's buffer
```

```
hello, world
```

```
.          period on a line by itself stops append mode
```

```
1,$p      print first through last lines to see what's there
```

```
hello, world
```

```
w temp.foo write the buffer to a file
```

```
13        editor says it wrote 13 bytes
```

```
^?        type the interrupt character
```


```
?         editor catches signal, prints question mark
```

```
caught SIGINT and so does the parent process
```

```
q         leave editor
```

```
caught SIGCHLD
```

# Foreground and Background Process Groups



- SIGINT is sent to all three foreground processes
- Shell ignores it
- a.out process and *ed* catch the signal
- a.out should ignore the signals and *ed* only should catch the signal

# Correct POSIX.1 Pmplementation of system()

```
int system(const char *cmdstring)    /* with appropriate signal handling */
{
    pid_t          pid;
    int            status;
    struct sigaction ignore, saveintr, savequit;
    sigset_t        chldmask, savemask;
    if (cmdstring == NULL) return(1); /* command processor with UNIX */
    ignore.sa_handler = SIG_IGN;      /* ignore SIGINT and SIGQUIT */
    sigemptyset(&ignore.sa_mask);
    ignore.sa_flags = 0;
    if (sigaction(SIGINT, &ignore, &saveintr) < 0) return(-1);
    if (sigaction(SIGQUIT, &ignore, &savequit) < 0) return(-1);
    sigemptyset(&chldmask);          /* now block SIGCHLD */
    sigaddset(&chldmask, SIGCHLD);
    if (sigprocmask(SIG_BLOCK, &chldmask, &savemask) < 0) return(-1);
```

# system() function Contd ...

```
if ((pid = fork()) < 0) {
    status = -1;    /* probably out of processes */
} else if (pid == 0) {    /* child */
    /* restore previous signal actions & reset signal mask */
    sigaction(SIGINT, &saveintr, NULL);
    sigaction(SIGQUIT, &savequit, NULL);
    sigprocmask(SIG_SETMASK, &savemask, NULL);

    execl("/bin/sh", "sh", "-c", cmdstring, (char *)0);
    _exit(127);    /* exec error */
} else {    /* parent */
    while (waitpid(pid, &status, 0) < 0)
        if (errno != EINTR) {
            status = -1; /* error other than EINTR from waitpid() */
            break;
        }
}

/* restore previous signal actions & reset signal mask */
if (sigaction(SIGINT, &saveintr, NULL) < 0)
    return(-1);
if (sigaction(SIGQUIT, &savequit, NULL) < 0)
    return(-1);
if (sigprocmask(SIG_SETMASK, &savemask, NULL) < 0)
    return(-1);

return(status);
```

# POSIX Implementation

- No signal is sent to caller when interrupt or quit character is typed
- When ed exits, SIGCHLD is sent to caller and it is blocked. Caller calls waitpid, collects termination status. POSIX.1 states that if wait or waitpid returns the status of a child process while SIGCHLD is pending, then SIGCHLD should not be delivered to the process unless the status of another child process is also available.
- If the caller has any pending SIGCHLD (of own children), that will be pending after waitpid and will be unblocked after sigprocmask.



# sleep ()

- unsigned int sleep(unsigned int seconds);  
Returns: 0 or number of unslept seconds
- This function causes the calling process to be suspended until either
  1. The amount of wall clock time specified by seconds has elapsed
    - Returns 0
  2. A signal is caught by the process and the signal handler returns
    - Returns number of unslept seconds
- if we do an alarm(10) and 3 wall clock seconds later do a sleep(5), what happens?
  - Sleep() will return in 5 seconds (assuming that some other signal isn't caught in that time), but will another SIGALRM be generated 2 seconds later?
  - These details depend on the implementation.

# sleep() - Simple, incomplete implementation

```
#include      <signal.h>
#include      <unistd.h>

static void
sig_alm(int signo)
{
    return; /* nothing to do, just return to wake up the pause */
}

unsigned int
sleep1(unsigned int nsecs)
{
    if (signal(SIGALRM, sig_alm) == SIG_ERR)
        return(nsecs);
    alarm(nsecs);           /* start the timer */
    pause();                /* next caught signal wakes us up */
    return( alarm(0) );     /* turn off timer, return unslept time */
}
```

# sleep() implementation - Problems

- If the caller has already an alarm set, that alarm is erased by the first call to alarm
- We have modified the disposition of SIGALRM, we should save the disposition and restore it
- Race condition between first call to alarm and pause

# Another (imperfect) implementation of sleep (SVR2)

```
static jmp_buf    env_alrm;

static void
sig_alrm(int signo)
{
    longjmp(env_alrm, 1);
}

unsigned int
sleep2(unsigned int nsecs)
{
    if (signal(SIGALRM, sig_alrm) == SIG_ERR)
        return(nsecs);
    if (setjmp(env_alrm) == 0) {
        alarm(nsecs);
        pause();
    }
    return( alarm(0) );
}
```

# Problem with implementation

- Race condition is eliminated
- Interaction with other signals
  - If SIGALRM interrupts some other signal handler, when `longjmp` is called, it aborts the other signal handler.

# Calling sleep2 from a program that catches other signals

```
unsigned int      sleep2(unsigned int);
static void      sig_int(int);

int main(void)
{
    unsigned int    unslept;
    if (signal(SIGINT, sig_int) == SIG_ERR) err_sys("signal(SIGINT) error");
    unslept = sleep2(5);
    printf("sleep2 returned: %u\n", unslept);
    exit(0);
}

static void sig_int(int signo)
{
    int i;
    volatile int    j;
    printf("\nsig_int starting\n");
    for (i = 0; i < 2000000; i++) j += i * i;
    printf("sig_int finished\n");
    return;
}
```



# sleep in Unix flavors

- Solaris 9 implements sleep using alarm. Manual page says that a previously scheduled alarm is properly handled.
  - For example, in the preceding scenario, before sleep returns, it will reschedule the alarm to happen 2 seconds later; sleep returns 0 in this case.
  - Also, if we do an alarm(6) and 3 wall clock seconds later do a sleep(5), the sleep returns in 3 seconds, not in 5 seconds. Here, the return value from sleep is 2.
- FreeBSD 5.2.1, Linux 2.4.22, and Mac OS X 10.3, the delay is provided by nanosleep, a function specified to be a high-resolution delay by the real-time extensions in the Single UNIX Specification. This function allows the implementation of sleep to be independent of signals.

# Posix.1 Implementation of sleep()

- Handles signals reliably, avoid race conditions between alarm and pause
- No effect on other signal handler that may be executing when SIGALRM is handled
- Do not handle interaction with previously set alarms.

# Posix.1 Implementation of sleep()

```
static void
sig_alm(int signo)
{
    /* nothing to do, just returning wakes up sigsuspend() */
}
```

```
unsigned int
sleep(unsigned int nsecs)
{
    struct sigaction    newact, oldact;
    sigset_t            newmask, oldmask, suspmask;
    unsigned int        unslept;

    /* set our handler, save previous information */
    newact.sa_handler = sig_alm;
    sigemptyset(&newact.sa_mask);
    newact.sa_flags = 0;
    sigaction(SIGALRM, &newact, &oldact);

    /* block SIGALRM and save current signal mask */
    sigemptyset(&newmask);
    sigaddset(&newmask, SIGALRM);
    sigprocmask(SIG_BLOCK, &newmask, &oldmask);

    alarm(nsecs);
```

# Implementation of sleep()

```
suspmask = oldmask;
sigdelset(&suspmask, SIGALRM);    /* make sure SIGALRM isn't blocked */
sigsuspend(&suspmask);            /* wait for any signal to be caught */

/* some signal has been caught,  SIGALRM is now blocked */

unslept = alarm(0);
sigaction(SIGALRM, &oldact, NULL); /* reset previous action */

/* reset signal mask, which unblocks SIGALRM */
sigprocmask(SIG_SETMASK, &oldmask, NULL);
return(unslept);
}
```

# Values with signal masks

- oldmask - initial when process starts
- newmask - block SIGALRM // 1. any previously set alarm expires, should not get confused. 2. if Alarm expires before pause, SIGALRM is blocked
- suspmask- oldmask – SIGALRM// ensure SIGALRM not blocked, some signal will be caught
- //if alarm expires before pause, blocked SIGALRM is unblocked by sigsuspend call, delivers SIGALRM and breaks pause
- newmask - block SIGALRM //pause broke due to some other signal, alarm expires, block SIGALRM
- Oldmask - //unblocks SIGALRM